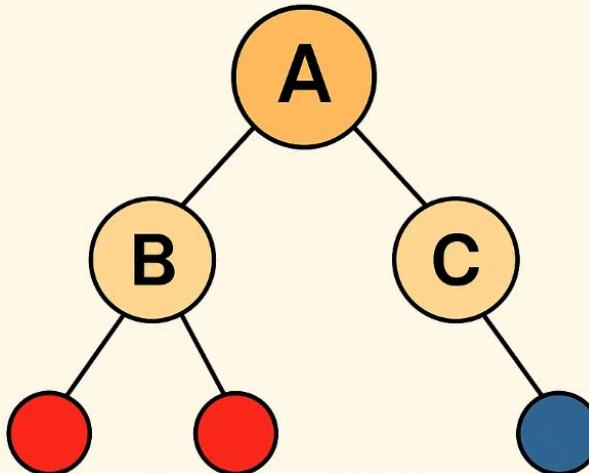


---

# CART for Classification

Mingrui Chen  
Muxin Fu  
Jingmin Xu  
Yixiao Zhang

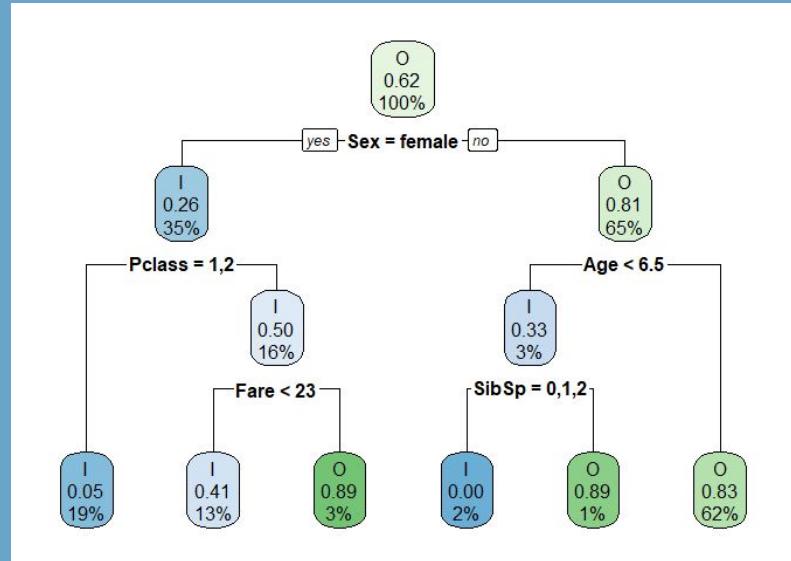
## Classification Trees (CART)



# Meet CART



- CART (Classification and Regression Trees) is a foundational supervised learning algorithm used for classification and regression tasks.
- It works by recursively partitioning the feature space into regions that are increasingly “pure” with respect to the target variable.
- The algorithm selects splits based on impurity measures (e.g., Gini index, entropy) to build an interpretable tree structure.
- CART remains widely used today due to its interpretability, flexibility, and role in ensemble models (Random Forests, Gradient Boosting).



Reference: <https://machinelearning-basics.com/cart-algorithm-and-everything-you-need-to-know-about-decision-trees-1-2/>

---

# Advantages & Disadvantages

## Advantages:

- Allow split on **continuous** features (!)
  - Threshold-based Split
- Highly interpretable
- Captures nonlinear relationships
- Minimal preprocessing required

## Disadvantages:

- Prone to overfitting
- Unstable
- Limited predictive power compared with ensembles

---

# **Math behind the ML algorithm**

## **Part 1: Representation**

# Part 1: Representation

---

- a. Domain Set  $\mathcal{X} = \mathbb{R}^n, \quad x_i = (x_{i1}, x_{i2}, \dots, x_{in}) \in \mathcal{X}.$
- b. Label Set  $\mathcal{Y} = \{0, 1, \dots, K - 1\}.$
- c. Training Data  $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^N, \quad x_i \in \mathcal{X}, \quad y_i \in \mathcal{Y}.$
- d. Learner's Output  $\mathcal{H} = \{ h : \mathcal{X} \rightarrow \mathcal{Y} \mid h \text{ is a binary decision tree with depth } \leq T_{\max} \}.$

Model as a Hierarchical Partition

- Each split defined by a feature  $j$  and threshold  $t$ :

$$\begin{aligned} x_j \leq t &\Rightarrow \text{left child} \\ x_j > t &\Rightarrow \text{right child} \end{aligned}$$

---

# **Math behind the ML algorithm**

## **Part 2: Loss**

## Part 2.1: Loss - Impurity Measure

---

Gini:

$$G_i = 1 - \sum_{k=1}^K p_{i,k}^2.$$

Entropy:

$$H_i = - \sum_{k=1}^K p_{i,k} \log p_{i,k},$$

- a. Node  $i$  containing a subset of samples:

$$S_i = \{(x_j, y_j)\}_{j \in \mathcal{I}_i}, \quad N_i = |S_i|.$$

- b. The number of samples in node  $i$  that belongs to class  $k$ :

$$n_{i,k} = \sum_{j \in \mathcal{I}_i} \mathbf{1}(y_j = k).$$

- c. Class proportion of class  $k$  in node  $i$ :

$$p_{i,k} = \frac{n_{i,k}}{N_i}, \quad k = 1, \dots, K.$$

$$\sum_{k=1}^K p_{i,k} = 1, \text{ and } p_{i,k} \geq 0 \quad \text{for } k = 1, \dots, K.$$

## Part 2.2: Loss - Split Loss

---

CART chooses the split that minimizes impurity of child nodes.

- The general **objective** therefore is:

$$L(S_i, \theta) = \frac{N_i^{\text{left}}}{N_i} C(S_i^{\text{left}}(\theta)) + \frac{N_i^{\text{right}}}{N_i} C(S_i^{\text{right}}(\theta)).$$

- At node  $i$ , the dataset  $S_i$  is partitioned into a left subset and right subset

$$S_i^{\text{left}}(\theta) = \{(x_j, y_j) \in S_i \mid x_{j,f} \leq t\},$$

$$S_i^{\text{right}}(\theta) = S_i \setminus S_i^{\text{left}}(\theta),$$

$$N_i^{\text{left}} = |S_i^{\text{left}}(\theta)|, \quad N_i^{\text{right}} = |S_i^{\text{right}}(\theta)|.$$

- It calculated a weighted child impurity
- $C(\cdot)$  is the impurity measure function

---

# Math behind the ML algorithm

## Part 3: Optimization

## Part 3.1: Greedy Search

---

- CART optimize to select the parameters that minimises the impurity at each split:

$$\theta^* = \arg \min_{\theta} L(S_i, \theta).$$

- **Greedy Optimization**

- For each node,
    - Evaluate every feature  $j$
    - Evaluate every candidate threshold  $t$
    - Compute impurity of left/right partitions
    - Choose the best pair
  - For example, at node  $i$ , we have impurity measure:

$$L(S_i, \theta) = \frac{N_i^{\text{left}}}{N_i} C(S_i^{\text{left}}(\theta)) + \frac{N_i^{\text{right}}}{N_i} C(S_i^{\text{right}}(\theta)).$$

## Part 3.2: Threshold & Stopping

---

- **Threshold candidate generation:**

- For sorted feature column, take **midpoints** as threshold candidates
  - For feature column  $v_1 < v_2 < \dots < v_m$ , we have:

$$\mathcal{T}_j = \left\{ t_k = \frac{v_k + v_{k+1}}{2} \mid k = 1, \dots, m-1 \right\}$$

- Example:

- For feature j, it has values: {1, 1, 2, 2, 4}
    - Unique sorted : {1,2,4}
    - Candidate thresholds: t1 = 1.5, t2 = 3.0

- **Stopping Criterion:**

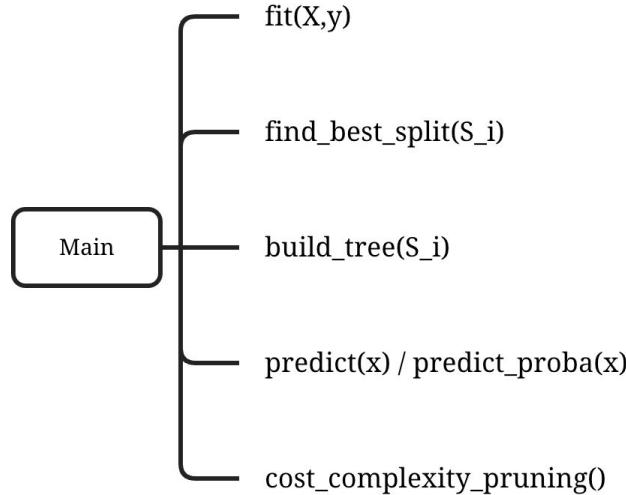
$$\text{Stop}(S_i) = \left( \underbrace{c(S_i) = 0}_{\text{pure node}} \right) \vee \left( \underbrace{|S_i| < \text{min\_samples\_split}}_{\text{too few samples}} \right) \vee \left( \underbrace{\text{depth}(S_i) \geq \text{max\_depth}}_{\text{depth limit}} \right) \vee \left( \underbrace{\Theta_i = \emptyset}_{\text{no valid split}} \right)$$

---

# Numerical Techniques

# Part 1: Numerical Structure

---



## Part 2.1: find the best split

```
FUNCTION FIND_BEST_SPLIT(S_i):  
    best_feature ← None  
    best_threshold ← None  
    best_loss ← +∞  
  
    FOR each feature j in 1..d:      → Evaluate each feature as a candidate  
        x_j ← column j of S_i  
        v ← sorted unique values of x_j  
  
        IF length(v) == 1:  
            CONTINUE // no valid split  
  
        T_j ← midpoints of consecutive values in v → Generate threshold midpoints  
  
        FOR each t in T_j:  
            S_left ← samples with x_j ≤ t  
            S_right ← samples with x_j > t  
  
            IF S_left empty OR S_right empty:  
                CONTINUE  
  
            imp_left ← impurity(S_left)  
            imp_right ← impurity(S_right)  
  
            loss ← (|S_left|/|S_i|)*imp_left + (|S_right|/|S_i|)*imp_right  
  
            IF loss < best_loss:  
                best_loss ← loss  
                best_feature ← j  
                best_threshold ← t  
  
    RETURN (best_feature, best_threshold)
```

- find the best feature-threshold pair  $(j^*, t^*)$  minimizing node loss.
- Return best\_feature, best\_threshold for each node

→ Evaluate each threshold, and update the best split

## Part 2.2: build\_tree recursively

```
FUNCTION BUILD_TREE(S_i, depth):  
    imp ← impurity(S_i)  
    n_samples ← number of samples in S_i  
  
    1. IF imp == 0:  
        # Node is pure – all samples belong to the same class  
        RETURN LeafNode(class_distribution(S_i))  
  
    2. IF n_samples < min_samples_split:  
        # Too few samples to reliably split  
        RETURN LeafNode(class_distribution(S_i))  
  
    3. IF depth ≥ max_depth:  
        # Depth limit reached  
        RETURN LeafNode(class_distribution(S_i))  
  
    4. IF no_valid_threshold_exists(S_i):  
        # All features have only one unique value, or all splits produce empty child nodes  
        RETURN LeafNode(class_distribution(S_i))  
  
(j*, t*) ← FIND_BEST_SPLIT(S_i) → Find best split parameters (feature & threshold)  
  
IF j* is None:  
    RETURN LeafNode(class_distribution(S_i))  
  
S_left ← samples in S_i where x_j* ≤ t*  
S_right ← samples in S_i where x_j* > t*  
  
LeftChild ← BUILD_TREE(S_left, depth + 1)  
RightChild ← BUILD_TREE(S_right, depth + 1)  
  
RETURN InternalNode(  
    feature = j*,  
    threshold = t*,  
    left = LeftChild,  
    right = RightChild  
)
```

- Recursively grow the CART tree
- using greedy splits and defined stopping rules.

→ Stopping Conditions

- Sklearn has more details constraints, we only implement the main controls and ignore some of it (like splitter, use *splitter = best*, assume *min\_impurity\_decrease=0.0* etc)

→ Split dataset, and grow recursively

## Part 2.2: cost\_complexity pruning

```
FUNCTION COST_COMPLEXITY_PRUNING_PATH(tree):

    # Step 1 - Compute subtree impurity and leaf count (bottom-up)
    FOR each node t (post-order traversal):
        IF t is a leaf:
            R_subtree[t] ← impurity(t) * sample_count(t)
            num_leaves[t] ← 1
        ELSE:
            R_subtree[t] ← R_subtree[left_child(t)] + R_subtree[right_child(t)]
            num_leaves[t] ← num_leaves[left_child(t)] + num_leaves[right_child(t)]

    # Step 2 - Compute the "cost of pruning"  $\alpha_t$  for each internal node
    FOR each internal node t:
        R_leaf ← impurity( t treated as a leaf ) * sample_count(t)
         $\alpha_t$  ← ( R_leaf - R_subtree[t] ) / ( num_leaves[t] - 1 )

    # Step 3 - Iteratively prune subtrees with the smallest  $\alpha_t$ 
     $\alpha$ _list ← [ ]
    impurity_list ← [ R_subtree[root] ]

    WHILE the tree can still be pruned:
         $\alpha_{\min}$  ← minimum  $\alpha_t$  over all remaining internal nodes
        prune all subtrees whose  $\alpha_t = \alpha_{\min}$  (replace with a leaf)
        update the total tree impurity
        append  $\alpha_{\min}$  to  $\alpha$ _list
        append updated impurity to impurity_list

    RETURN (  $\alpha$ _list, impurity_list )
```

- Larger alpha → Stronger penalty  
→ More pruning
- Default alpha = 0, full tree
- The main idea is what's the cost of each internal node, if we cut it
- Cost can be understood as the increase of impurity / number of removed leaves

---

# Sklearn Results

# 1. Impurity Measure

Testing all features (forcing sklearn to split on each):

Feature 0 (sepal length (cm)), threshold=5.4500:  
Left - Sklearn: 0.2374260355, Custom: 0.2374260355, Diff: 0.00e+00, Match: True  
Right - Sklearn: 0.5458142441, Custom: 0.5458142441, Diff: 0.00e+00, Match: True

Feature 1 (sepal width (cm)), threshold=3.3500:  
Left - Sklearn: 0.6251076827, Custom: 0.6251076827, Diff: 0.00e+00, Match: True  
Right - Sklearn: 0.2790357925, Custom: 0.2790357925, Diff: 0.00e+00, Match: True → Gini

Feature 2 (petal length (cm)), threshold=2.4500:  
Left - Sklearn: 0.0000000000, Custom: 0.0000000000, Diff: 0.00e+00, Match: True  
Right - Sklearn: 0.5000000000, Custom: 0.5000000000, Diff: 0.00e+00, Match: True

Feature 3 (petal width (cm)), threshold=0.8000:  
Left - Sklearn: 0.0000000000, Custom: 0.0000000000, Diff: 0.00e+00, Match: True  
Right - Sklearn: 0.5000000000, Custom: 0.5000000000, Diff: 0.00e+00, Match: True

Testing all features (forcing sklearn to split on each):

Feature 0 (sepal length (cm)), threshold=5.5500:  
Left - Sklearn: 0.8128223064, Custom: 0.8128223064, Diff: 5.55e-16, Match: True  
Right - Sklearn: 1.1670654490, Custom: 1.1670654490, Diff: 6.66e-16, Match: True

Feature 1 (sepal width (cm)), threshold=3.3500:  
Left - Sklearn: 1.4842073301, Custom: 1.4842073301, Diff: 4.44e-16, Match: True  
Right - Sklearn: 0.7448661738, Custom: 0.7448661738, Diff: 1.11e-16, Match: True → Entropy

Feature 2 (petal length (cm)), threshold=2.4500:  
Left - Sklearn: 0.0000000000, Custom: 0.0000000000, Diff: 0.00e+00, Match: True  
Right - Sklearn: 1.0000000000, Custom: 1.0000000000, Diff: 0.00e+00, Match: True

Feature 3 (petal width (cm)), threshold=0.8000:  
Left - Sklearn: 0.0000000000, Custom: 0.0000000000, Diff: 0.00e+00, Match: True  
Right - Sklearn: 1.0000000000, Custom: 1.0000000000, Diff: 0.00e+00, Match: True

- We have the **same impurity measure** result at each feature compared with sklearn
- This test was created to make sure floating precision (to check **tie-breaking** problem later on)

## 2. Performance: seed = 0

---

### Experimental Setup

- Dataset: Iris
- Train/Test Split: 70% / 30%,  
stratified, random\_state = 0
- Models Compared:
  - sklearn.tree.DecisionTreeClassifier
  - Our DecisionTreeCART implementation

### Dataset: IRIS (GINI)

---

```
Exact match: True
Same acc?: True
Prediction agreement: 100.00%
sklearn accuracy: 0.9556
our CART accuracy: 0.9556
SUCCESS: Exact match achieved on the Iris dataset.
```

### Dataset: IRIS (ENTROPY)

---

```
Exact match: True
Same acc?: True
Prediction agreement: 100.00%
sklearn accuracy: 0.9556
our CART accuracy: 0.9556
SUCCESS: Exact match achieved on the Iris dataset.
```

## 2. Performance: More

Seed	Our CART Acc	Sklearn Acc	Diff (Our-Sk)	#Mismatches	Notes
0	0.9778	0.9778	+0.0000	0	Predictions identical
1	0.9778	0.9778	+0.0000	0	Predictions identical
3	0.9333	0.9333	+0.0000	0	Predictions identical
5	0.9333	0.9333	+0.0000	0	Predictions identical
10	0.9778	1.0000	-0.0222	1	Our model misclassifies 1 sample
33	0.9556	0.9111	+0.0444	2	Sklearn misclassifies 2 samples

Seed	Our CART Acc	SKlaern Acc	Diff (Our-Sk)	# Mismatch	Notes
0	0.9556	0.9556	+0.0000	0	Predictions identical
1	0.9556	0.9778	-0.0222	1	idx=15, X=[6.8,2.8,4.8,1.4], y=1, our=2, sk=1
5	0.9333	0.9333	+0.0000	0	Predictions identical
10	0.9778	1.0000	-0.0222	1	idx=9, X=[4.9,2.4,3.3,1.0], y=1, our=2, sk=1

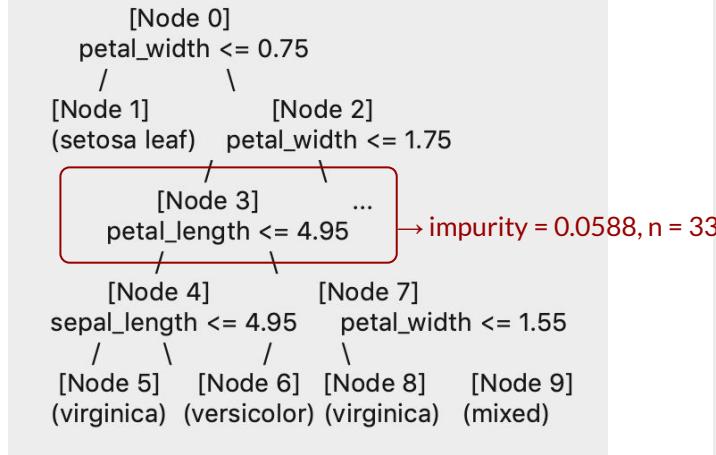
### Experimental Setup

- Dataset: Iris
- Train/Test Split: 70% / 30%
- Max\_depth = 5
- Min\_samples\_split = 2
- We reached almost the same prediction and split as what sklearn does
- Only with minor mismatch due

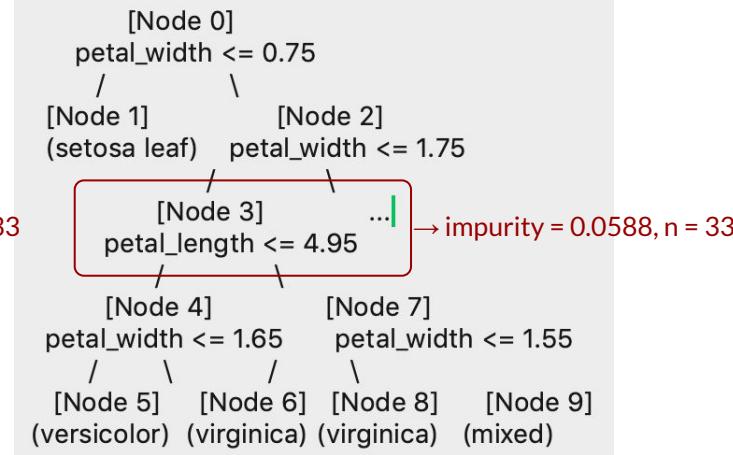
→ Gini

→ Entropy

### 3. Tie-breaking Condition



Ours



Sklearn

- We future explore the case when seed=10, why we have one mismatch
- We found that it happens at the tree node 3 -> Where our model select feature 2, and sklearn select feature 3 to split one
- However, this two features **have the same impurity measure** -> Tie breaking condition
- This may due to the difference in how sklearn implement randomness, and how they deal with the same gain features.

# Summary Slide

## What We Found Interesting

- CART's decision logic is intuitive yet elegant:  
It evaluates all features, sorts values, and picks the split with the lowest impurity.
- The model is highly interpretable:  
Each node's split can be explained visually and logically—rare among ML models.

## What Was Challenging

- Debugging tree mismatches: Comparing node structures, impurity calculations, and thresholds was surprisingly demanding.
-

# A little bit more..

- Random State Control:
  - In Sklearn documents, it said:
  - We implement it, permute feature every split, using numpy.random.RandomState
- But the source code, written in cython, it's:

```
self.rand_r_state = self.random_state.randint(0, RAND_R_MAX)
cdef intp_t n_samples = X.shape[0]
```

- This may lead to the different permutation order, and therefore the way how we encounter the features that have the same gain will be different
  - Since use strictly gain comparison, the first met one will be taken
- Also, float32 / float64 will have different precision.

## Source Code:

[https://github.com/scikit-learn/scikit-learn/blob/1eb422d6c5f46a98a318f341de3e4709f9521bfe/sklearn/tree/\\_splitter.pyx](https://github.com/scikit-learn/scikit-learn/blob/1eb422d6c5f46a98a318f341de3e4709f9521bfe/sklearn/tree/_splitter.pyx)

### random\_state : int, RandomState instance or None, default=None

Controls the randomness of the estimator. The features are always randomly permuted at each split, even if `splitter` is set to "best". When `max_features < n_features`, the algorithm will select `max_features` at random at each split before finding the best split among them. But the best found split may vary across different runs, even if `max_features=n_features`. That is the case, if the improvement of the criterion is identical for several splits and one split has to be selected at random. To obtain a deterministic behaviour during fitting, `random_state` has to be fixed to an integer. See [Glossary](#) for details.

### random\_state

Whenever randomization is part of a Scikit-learn algorithm, a `random_state` parameter may be provided to control the random number generator used. Note that the mere presence of `random_state` doesn't mean that randomization is always used, as it may be dependent on another parameter, e.g. `shuffle`, being set.

The passed value will have an effect on the reproducibility of the results returned by the function ([fit](#), [split](#), or any other function like [k\\_means](#)). `random_state`'s value may be:

#### None (default)

Use the global random state instance from [numpy.random](#). Calling the function multiple times will reuse the same instance, and will produce different results.

#### An integer

Use a new random number generator seeded by the given integer. Using an int will produce the same results across different calls. However, it may be worthwhile checking that your results are stable across a number of different distinct random seeds. Popular integer random seeds are 0 and 42. Integer values must be in the range [0, 2\*\*32 - 1].

#### A [numpy.random.RandomState](#) instance

Use the provided `random_state`, only affecting other users of that same random state instance. Calling the function multiple times will reuse the same instance, and will produce different results.

---

Thank you!

Q & A